

# Homework 4:

## Reinforcement Learning

Part I. Implementation (-5 if not explain in detail):

- Please screenshot your code snippets of Part 1 ~ Part 3, and explain your implementation
- Part 1, taxi

```
1 def choose_action(self, state):
2     """
3     Choose the best action with given state and epsilon.
4
5     Parameters:
6         state: A representation of the current state of the environment.
7         epsilon: Determines the explore/exploit rate of the agent.
8
9     Returns:
10        action: The action to be evaluated.
11    """
12    # Begin your code
13    # TODO
14    # raise NotImplementedError("Not implemented yet.")
15    # Choose the best action according to the qtable and epsilon
16    if np.random.uniform(0, 1) < self.epsilon or np.sum(self.qtable[state]) == 0:
17        # Explore
18        action = self.env.action_space.sample()
19    else:
20        # Exploit
21        action = np.argmax(self.qtable[state])
22    return action
23    # End your code
24
25 def learn(self, state, action, reward, next_state, done):
26     """
27     Calculate the new q-value base on the reward and state transformation observed after taking the action.
28
29     Parameters:
30         state: The state of the environment before taking the action.
31         action: The executed action.
32         reward: Obtained from the environment after taking the action.
33         next_state: The state of the environment after taking the action.
34         done: A boolean indicates whether the episode is done.
35
36     Returns:
37         None (Don't need to return anything)
38    """
39    # Begin your code
40    # TODO
41    # raise NotImplementedError("Not implemented yet.")
42    # Calculate the new Q-value based on the Q-learning formula
43    q_next = np.max(self.qtable[next_state]) if not done else 0
44    q_current = self.qtable[state, action]
45
46    # Update the Q-table with the new Q-value
47    q_newValue = (1 - self.learning_rate) * q_current + self.learning_rate \
48        * (reward + self.gamma * q_next)
49
50    # Record the Q-value
51    self.qtable[state, action] = q_newValue
52    # End your code
53    np.save("./Tables/taxi_table.npy", self.qtable)
54
```

```

54
55 def check_max_Q(self, state):
56     """
57     - Implement the function calculating the max Q value of given state.
58     - Check the max Q value of initial state
59
60     Parameter:
61         state: the state to be check.
62     Return:
63         max_q: the max Q value of given state
64     """
65     # Begin your code
66     # TODO
67     # raise NotImplementedError("Not implemented yet.")
68     # Find the max Q-value of the given state
69     max_q = np.max(self.qtable[state])
70     return max_q
71     # End your code"

```

## ● Part 2, cartpole

```

1 def init_bins(self, lower_bound, upper_bound, num_bins):
2     """
3     Slice the interval into #num_bins parts.
4     Parameters:
5         lower_bound: The lower bound of the interval.
6         upper_bound: The upper bound of the interval.
7         num_bins: Number of parts to be sliced.
8     Returns:
9         a numpy array of #num_bins - 1 quantiles.
10    Example:
11        Let's say that we want to slice [0, 10] into five parts,
12        that means we need 4 quantiles that divide [0, 10].
13        Thus the return of init_bins(0, 10, 5) should be [2. 4. 6. 8.].
14    Hints:
15        1. This can be done with a numpy function.
16    """
17    # Begin your code
18    # TODO
19    # raise NotImplementedError("Not implemented yet.")
20    # Slice the interval into #num_bins parts
21    array = np.linspace(lower_bound, upper_bound, num_bins + 1)
22    return array[1:-1]
23    # End your code
24
25 def discretize_value(self, value, bins):
26     """
27     Discretize the value with given bins.
28     Parameters:
29         value: The value to be discretized.
30         bins: A numpy array of quantiles
31     returns:
32         The discretized value.
33    Example:
34        With given bins [2. 4. 6. 8.] and "5" being the value we're going to discretize.
35        The return value of discretize_value(5, [2. 4. 6. 8.]) should be 2, since 4 <= 5 < 6 where [4, 6) is the 3rd bin.
36    Hints:
37        1. This can be done with a numpy function.
38    """
39    # Begin your code
40    # TODO
41    # raise NotImplementedError("Not implemented yet.")
42    # Discretize the value with given bins
43    return np.digitize(value, bins)
44    # End your code
45
46 def discretize_observation(self, observation):
47     """
48     Discretize the observation which we observed from a continuous state space.
49     Parameters:
50         observation: The observation to be discretized, which is a list of 4 features:
51             1. cart position.
52             2. cart velocity.
53             3. pole angle.
54             4. tip velocity.
55     Returns:
56         state: A list of 4 discretized features which represents the state.
57     Hints:
58         1. All 4 features are in continuous space.
59         2. You need to implement discretize_value() and init_bins() first
60         3. You might find something useful in Agent.__init__()
61     """
62     # Begin your code
63

```

```

63     # Begin your code
64     # TODO
65     # raise NotImplementedError("Not implemented yet.")
66     state = []
67     for index in range(4):
68         # Discretize the observation
69         state.append( self.discretize_value(observation[index], self.bins[index]) )
70     return state
71     # End your code
72
73 def choose_action(self, state):
74     """
75     Choose the best action with given state and epsilon.
76     Parameters:
77         state: A representation of the current state of the environment.
78         epsilon: Determines the explore/exploit rate of the agent.
79     Returns:
80         action: The action to be evaluated.
81     """
82     # Begin your code
83     # TODO
84     # raise NotImplementedError("Not implemented yet.")
85     # Choose the best action according to the qtable and epsilon
86     if np.random.uniform(0, 1) < self.epsilon or np.sum(self.qtable[state]) == 0:
87         # Explore
88         action = self.env.action_space.sample()
89     else:
90         # Exploit
91         action = np.argmax(self.qtable[tuple(state)])
92     return action
93     # End your code
94
95 def learn(self, state, action, reward, next_state, done):
96     """
97     Calculate the new q-value base on the reward and state transformation observed after taking the action.
98     Parameters:
99         state: The state of the environment before taking the action.
100         action: The executed action.
101         reward: Obtained from the environment after taking the action.
102         next_state: The state of the environment after taking the action.
103         done: A boolean indicates whether the episode is done.
104     Returns:
105         None (Don't need to return anything)
106     """
107     # Begin your code
108     # TODO
109     # raise NotImplementedError("Not implemented yet.")
110     # Calculate the new Q-value based on the Q-learning formula
111     q_next = np.max(self.qtable[tuple(next_state)]) if not done else 0
112     q_current = self.qtable[tuple(state)][action]
113
114     # Update the Q-table with the new Q-value
115     q_newValue = (1 - self.learning_rate) * q_current + self.learning_rate \
116                 * (reward + self.gamma * q_next)
117
118     # Record the Q-value
119     self.qtable[tuple(state)][action] = q_newValue
120     # End your code
121     np.save("../Tables/cartpole_table.npy", self.qtable)
122
123 def check_max_Q(self):
124     """
125     - Implement the function calculating the max Q value of initial state(self.env.reset()).
126     - Check the max Q value of initial state
127     Parameter:
128         self: the agent itself.
129         (Don't pass additional parameters to the function.)
130         (All you need have been initialized in the constructor.)
131     Return:
132         max_q: the max Q value of initial state(self.env.reset())
133     """
134     # Begin your code
135     # TODO
136     # raise NotImplementedError("Not implemented yet.")
137     state = self.discretize_observation(self.env.reset())
138     max_q = np.max(self.qtable[tuple(state)])
139     return max_q
140     # End your code

```


## ● Part 3, DQN

```
1 def learn(self):
2     """
3     - Implement the learning function.
4     - Here are the hints to implement.
5     Steps:
6     -----
7     1. Update target net by current net every 100 times. (we have done this for you)
8     2. Sample trajectories of batch size from the replay buffer.
9     3. Forward the data to the evaluate net and the target net.
10    4. Compute the loss with MSE.
11    5. Zero-out the gradients.
12    6. Backpropagation.
13    7. Optimize the loss function.
14    -----
15    Parameters:
16        self: the agent itself.
17        (Don't pass additional parameters to the function.)
18        (All you need have been initialized in the constructor.)
19    Returns:
20        None (Don't need to return anything)
21    """
22    if self.count % 100 == 0:
23        self.target_net.load_state_dict(self.evaluate_net.state_dict())
24
25    # Begin your code
26    # TODO
27    # raise NotImplementedError("Not implemented yet.")
28    # Sample trajectories of batch size from the replay buffer.
29    observations, actions, rewards, next_observations, done = self.buffer.sample(self.batch_size)
30
31    # Forward the data to the evaluate net and the target net.
32    observations = torch.FloatTensor(np.array(observations))
33    actions = torch.LongTensor(actions)
34    rewards = torch.FloatTensor(rewards)
35    next_observations = torch.FloatTensor(np.array(next_observations))
36    done = torch.BoolTensor(done)
37
38    # Compute the loss with MSE.
39    evaluate = self.evaluate_net(observations).gather(1, actions.reshape(self.batch_size, 1))
40    nextMax = self.target_net(next_observations).detach()
41    target = rewards.reshape(self.batch_size, 1) + self.gamma * nextMax.max(1)[0].view(self.batch_size, 1) * (~done).reshape(self.batch_size, 1)
42
43    # Zero-out the gradients.
44    MSE = nn.MSELoss()
45    loss = MSE(evaluate, target)
46    self.optimizer.zero_grad()
47    loss.backward()
48    self.optimizer.step()
49    # End your code
50    torch.save(self.target_net.state_dict(), "./Tables/DQN.pt")
51
52 def choose_action(self, state):
53     """
54     - Implement the action-choosing function.
55     - Choose the best action with given state and epsilon
56     Parameters:
57         self: the agent itself.
58         state: the current state of the environment.
59         (Don't pass additional parameters to the function.)
60         (All you need have been initialized in the constructor.)
61     Returns:
62         action: the chosen action.
63     """
64    with torch.no_grad():
65        # Begin your code
66        # TODO
67        # raise NotImplementedError("Not implemented yet.")
68        if random.uniform(0,1) < self.epsilon:
69            action = self.env.action_space.sample()
70        else:
71            action = torch.argmax(self.evaluate_net.forward(torch.FloatTensor(state))).item()
72        # End your code
73    return action
74
75 def check_max_q(self):
76     """
77     - Implement the function calculating the max Q value of initial state(self.env.reset()).
78     - Check the max Q value of initial state
79     Parameter:
80         self: the agent itself.
81         (Don't pass additional parameters to the function.)
82         (All you need have been initialized in the constructor.)
83     Return:
84         max_q: the max Q value of initial state(self.env.reset())
85     """
86    # Begin your code
87    # TODO
88    # raise NotImplementedError("Not implemented yet.")
89    max_q = torch.max(self.evaluate_net.forward(torch.FloatTensor(self.env.reset()))).item()
90    return max_q
91    # End your code
```


## Part II. Experiment Results:

- Please paste [taxi.png](#), [cartpole.png](#), [DQN.png](#) and [compare.png](#) here.


### 1. taxi.png:




### 2. cartpole.png



### 3. DQN.png



### 4. compare.png





### Part III. Question Answering (50%):

1. Calculate the optimal Q-value of a given state in Taxi-v3, and compare with the Q-value you learned (Please screenshot the result of the “`check_max_Q`” function to show the Q-value you learned). (10%)

$$Q_{\text{opt}}(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}(s')]$$

$Q_{\text{opt}} \approx 1.6226$ , approximately equals to the max Q we get.

```
(ai_hw4) alfonso@ubuntu:~/Git_workspace/Intro-to-AI-2023/HW4/AI_HW4$ python
3 taxi.py
#1 training progress
100%|████████████████████████████████████████| 3000/3000 [00:52<00:00, 56.96it/s]
#2 training progress
100%|████████████████████████████████████████| 3000/3000 [00:28<00:00, 105.46it/s]
#3 training progress
100%|████████████████████████████████████████| 3000/3000 [00:09<00:00, 314.10it/s]
#4 training progress
100%|████████████████████████████████████████| 3000/3000 [00:29<00:00, 101.72it/s]
#5 training progress
100%|████████████████████████████████████████| 3000/3000 [00:20<00:00, 146.03it/s]
average reward: 7.54
Initail state:
taxi at (2, 2), passenger at Y, destination at R
max Q:1.6226146699999995
```

2. Calculate the max Q-value of the initial state in CartPole-v0, and compare with the Q-value you learned. (Please screenshot the result of the “`check_max_Q`” function to show the Q-value you learned) (10%)

$$\frac{1}{1 - \gamma} = \frac{1}{1 - 0.97} = 33.3333$$

```
(ai_hw4) alfonso@ubuntu:~/Git_workspace/Intro-to-AI-2023/HW4/AI_HW4$ python
3 cartpole.py
#1 training progress
100%|████████████████████████████████████████| 3000/3000 [02:06<00:00, 23.63it/s]
#2 training progress
100%|████████████████████████████████████████| 3000/3000 [02:00<00:00, 24.84it/s]
#3 training progress
100%|████████████████████████████████████████| 3000/3000 [02:03<00:00, 24.39it/s]
#4 training progress
100%|████████████████████████████████████████| 3000/3000 [02:25<00:00, 20.67it/s]
#5 training progress
100%|████████████████████████████████████████| 3000/3000 [01:57<00:00, 25.54it/s]
average reward: 172.96
max Q:31.094193147630833
```

```
(ai_hw4) alfonso@ubuntu:~/Git_workspace/Intro-to-AI-2023/HW4/AI_HW4$ python
3 DQN.py
#1 training progress
100%|████████████████████████████████████████| 1000/1000 [02:01<00:00, 8.26it/s]
#2 training progress
100%|████████████████████████████████████████| 1000/1000 [02:26<00:00, 6.82it/s]
#3 training progress
100%|████████████████████████████████████████| 1000/1000 [02:34<00:00, 6.46it/s]
#4 training progress
100%|████████████████████████████████████████| 1000/1000 [02:24<00:00, 6.92it/s]
#5 training progress
100%|████████████████████████████████████████| 1000/1000 [08:44<00:00, 1.91it/s]
reward: 200.0
max Q:0.0336914137005806
```

3.

- a. Why do we need to discretize the observation in Part 2? (3%)

A: Discretize the observation would make the algorithm more easier to converge. In this case, the dimension of the state space is very high, and using continuous observations would make the algorithm difficult to train.

- b. How do you expect the performance will be if we increase “num\_bins” ? (3%)

A: Better, because by increasing the number of bins, we can increase the granularity of the discretization and make the state space more fine-grained. This can make the policy more precise and enable better policies to be learned.

- c. Is there any concern if we increase “num\_bins” ? (3%)

A: Increasing the number of bins will increase the dimension of the state space, which will make training time longer and require more memory to store the Q-table.

4. Which model (DQN, discretized Q learning) performs better in Cartpole-v0, and what are the reasons? (5%)

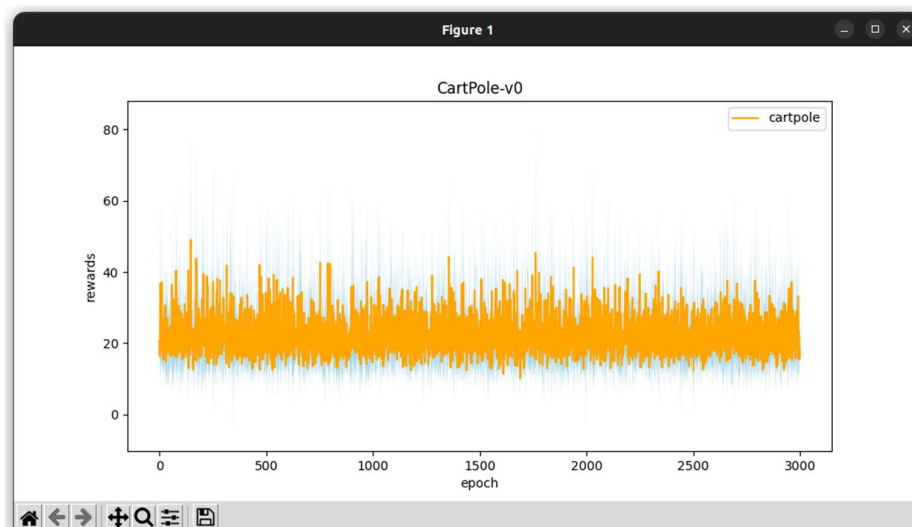
A: DQN, since CartPole-v0 has a continuous state space, which makes it challenging to create a discrete Q-table for Q-learning. Discretizing the state space can lead to a loss of information and reduced performance. On the other hand, DQN uses a neural network to approximate the Q-values, which can handle continuous state spaces more effectively.

5.

- a. What is the purpose of using the epsilon greedy algorithm while choosing an action? (3%)

The main purpose of the epsilon greedy algorithm is to choose between exploration and exploitation when the agent has none or limited knowledge about the environment.

- b. What will happen, if we don't use the epsilon greedy algorithm in the CartPole-v0 environment? (3%)





The result may look like the plot above, the agent is not going to explore the unknown environment.

- c. Is it possible to achieve the same performance without the epsilon greedy algorithm in the CartPole-v0 environment? Why or Why not? (3%)

Yes, softmax may reach same performance.

- d. Why don't we need the epsilon greedy algorithm during the testing section? (3%)

Because the agent has already known the environment.

6. Why does `"with torch.no_grad(): "` do inside the `"choose_action"` function in DQN? (4%)

A: It is used to disable gradient computation during the evaluation of the Q-values of the current state. During the testing phase, we only want to evaluate the learned policy and not update the weights of the network. Disabling gradient computation using `"with torch.no_grad():"` saves memory and computation time by preventing the calculation of gradients.